CSE305 Concurrent Programming: N-Body simulation project

Martin Lau, Oscar Peyron, Ziyue Qiu
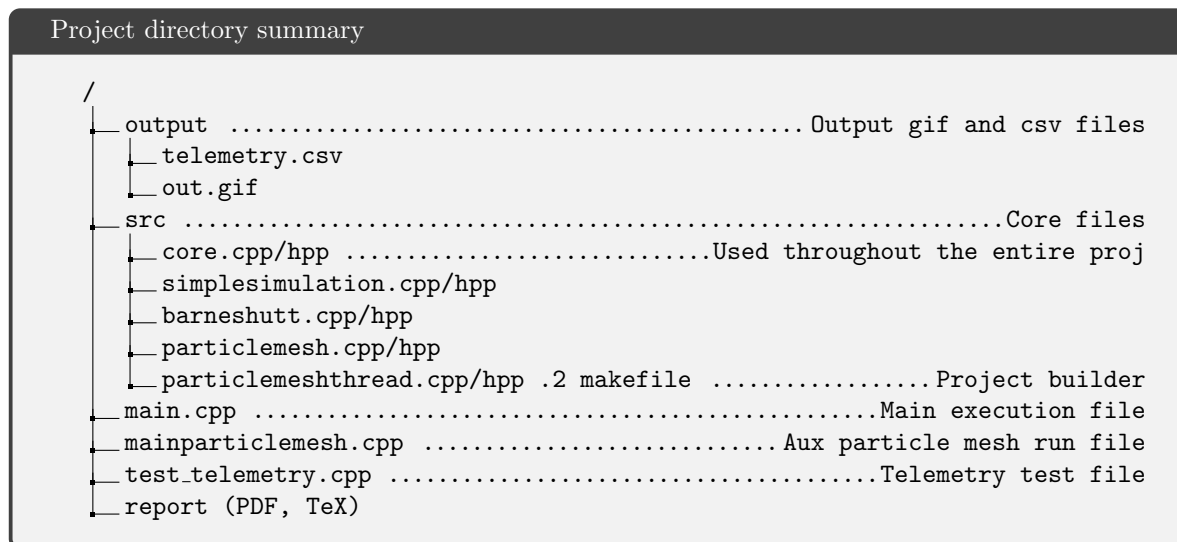
May 31, 2025

1 Introduction

This document outlines the N-Body simulation project for CSE305, which includes how we approached the problem, our project structure, code breakdown and strategies, and encountered difficulties.

To get started quickly, try typing `make nbody` in the project directory, and generate a basic gif file. For information on how to execute the code, check our README file.

For this project, our aim is to be able to simulate systems of bodies with forces interacting with one another in 2D (such as orbiting planets around the solar system with gravitational forces, or with particles interacting with each other with Coulomb forces). This therefore includes two sections:

1. Simulation of the N bodies and their evolving states.

2. Visualization of recorded telemetries into an animated output.

Below is a general view of our project repository file structure.

```
Project directory summary

/
├── output ............................................ Output gif and csv files
│   ├── telemetry.csv
│   └── out.gif
├── src ................................................................... Core files
│   ├── core.cpp/hpp ............................... Used throughout the entire proj
│   ├── simplesimulation.cpp/hpp
│   ├── barneshutt.cpp/hpp
│   ├── particlemesh.cpp/hpp
│   └── particlemeshthread.cpp/hpp .2 makefile .................. Project builder
├── main.cpp ................................................. Main execution file
├── mainparticlemesh.cpp ............................ Aux particle mesh run file
├── test_telemetry.cpp ........................................ Telemetry test file
└── report (PDF, TeX)
```

The work has been split as follows. The arrangement is not defining, as we like to help each other out in different parts. Our files have still been mostly separated to properly separate who worked on what functionalities.

1. Martin handles general project structure, core classes, visualization, and the naive and its optimized algorithm implementation.

2. Ziyue handles the Barnes Hutt algorithm.

3. Oscar handles the Particle Mesh algorithm implementation (simple, thread based and cuda implementation)

2 Core components

2.1 Main elements

There are 3 primary classes defined used throughout our project: `Vector, Body, System`. These are defined in `core.cpp/hpp`.

1. `Vector` holds information on a pair of numbers, and also allows for operations with other vectors/scalars (as opposed to using `std::pair`).

2. `Body` holds information on a given particle or body, including its mass, coordinates, velocity, and acceleration. It also contains an update method which updates its position and velocity based on acceleration.

3. `System` stores a collection of bodies and its recorded telemetry from the simulations we are going to do. It also contains the visualization function, which takes its recorded telemetry and outputs an animated file.

*Note: This organization is heavily inspired from one assignment from CSE306 Computer Graphics.*

Elaborating more on the telemetry stored within the `System` class, this is stored as a vector of vector of `Vector`. What a mouthful! Our simulations creates different steps/frames. Each step/frame contains the positions of all bodies inside the system (this is a vector of `Vector`). To have the entire telemetry, we have a vector of frames, or the aforementioned data structure for the telemetry.

2.2 Visualization code

To visualize the code, we opt to create a gif animation after the telemetry is recorded, using the `ImageMagick/Magick++` libraries. The bulk of the visualization code is located in `core.cpp` as a method for the `System` class. As of present, the visualization function works in two steps. First, it creates the frames for the animation, then writes the frames to a gif file. The frame creation is relatively quick, and most of the visualization time is actually spent in one line (`writeImages(frames.begin(), frames.end(), name)`).

We will look into further solutions to try and decrease the time taken to create the visualization gifs, like reducing image quality, or the number of frames.

For testing purposes, there is also a visualizer in `visualizer.py` which creates a animation using Python's `matplotlib` library, working significantly faster, allowing us to test the correctness of our telemetries.

3 Naive algorithm

So far, we have a basic implementation of the naive algorithm without multi threading (most of my time was taken bugfixing core functionalities and getting visualization to work).

Algorithm 1 Naive simulation outline

Require: System of bodies with masses m_i, initial positions $\vec{r}_i$, velocities $\vec{v}_i$
Require: Time step Δt, number of steps N
  telemetry $\leftarrow \emptyset$
  telemetry.append(initial positions)
  for step $\leftarrow 0$ to $N - 1$ **do**
    for each body i **do**
      $\vec{a}_i \leftarrow \vec{0}$ $\triangleright$ Reset accelerations
    end for
    for $i \leftarrow 0$ to $n - 1$ **do**
      for $j \leftarrow i + 1$ to $n - 1$ **do**
        $\vec{F}_{ij} \leftarrow$ ComputeGravitationalForce($body_i$, $body_j$)
        $\vec{a}_i \leftarrow \vec{a}_i + \vec{F}_{ij}/m_i$
        $\vec{a}_j \leftarrow \vec{a}_j - \vec{F}_{ij}/m_j$
      end for
    end for
    for each body i **do**
      $\vec{v}_i \leftarrow \vec{v}_i + \vec{a}_i \Delta t$ $\triangleright$ Update velocity
      $\vec{r}_i \leftarrow \vec{r}_i + \vec{v}_i \Delta t$ $\triangleright$ Update position
    end for
    telemetry.append(current positions)
  end for
Ensure: Position history for all bodies stored in telemetry

Other aspects (parallelizing the update steps, parallelizing forces computations, avoiding race conditions) will be implemented later on.

4 Particle Mesh based algorithm

I have implemented a sequential Particle mesh algorithm, as well as a thread based. Here below is the pseudo-code for the sequential implementation as well as for the thread-based implementation. For both implementations the assignment of the mass for each grid was done using a method called the Nearest Grid Point method. This means that the mass was assigned to the grid in which the particle is situated. Other methods for grid-assignment are known such as the "Cloud-In-Cell" or the "Triangle Shaped Cloud" but showed suboptimal results in parallel implementations.

Algorithm 2 Particle-Mesh Simulation, using Nearest-Grid-Point (NGP)

Require: System *universe*, time step Δt, grid size N, spatial extent R

1: telemetry $\leftarrow \emptyset$
2: telemetry.append(initial positions)
3: boundaries $\leftarrow [-R, R]$
4: Compute cell size $h \leftarrow \frac{2R}{N}$
5: Initialize mass density grid $M[G][G] \leftarrow 0$
6: Initialize potential grid $\Phi[G][G] \leftarrow 0$
7: Initialize FFTW input/output arrays and plans
8: **for** each time step $s = 1$ to S **do**
9:     Clear mass density grid M
10:     **for** each body b in *universe* **do**
11:         $i \leftarrow \left\lfloor \frac{x_b - \min_x}{h} \right\rfloor$
12:         $j \leftarrow \left\lfloor \frac{y_b - \min_y}{h} \right\rfloor$
13:         **if** (i, j) in bounds **then**
14:             $M[i][j] \leftarrow M[i][j] + \text{body.mass}$
15:         **end if**
16:     **end for**
17:     Copy grid mass to FFTW input array
18:     Compute FFT of mass density using forward FFT
19:     **for** each (i, j) in frequency domain **do**
20:         Compute wave numbers (k_x, k_y)
21:         Compute $k^2 \leftarrow k_x^2 + k_y^2$
22:         **if** $k^2 > 0$ **then**
23:             Multiply by $-\frac{G}{k^2}$ in frequency domain
24:         **else**
25:             Set value to zero
26:         **end if**
27:     **end for**
28:     Compute inverse FFT to obtain gravitational potential
29:     Normalize inverse FFT result and store in potential grid Φ
30:     **for** each body b in *universe* **do**
31:         $i \leftarrow \left\lfloor \frac{x_b - \min_x}{h} \right\rfloor$
32:         $j \leftarrow \left\lfloor \frac{y_b - \min_y}{h} \right\rfloor$
33:         **if** (i, j) is valid and not at boundary **then**
34:             Compute force via central difference of potential:
35:             $\vec{a}_i \leftarrow - \left(\frac{\Phi[i+1][j] - \Phi[i-1][j]}{2h}, \frac{\Phi[i][j+1] - \Phi[i][j-1]}{2h} \right)$
36:         **else**
37:             $\vec{a}_i \leftarrow (0, 0)$
38:         **end if**
39:     **end for**
40:     **for** each body i **do**
41:         $\vec{v}_i \leftarrow \vec{v}_i + \vec{a}_i \Delta t$ ▷ Update velocity
42:         $\vec{r}_i \leftarrow \vec{r}_i + \vec{v}_i \Delta t$ ▷ Update position
43:     **end for**
44:     telemetry.append(current positions)
45: **end for**
46: Cleanup FFTW plans and memory

Algorithm 3 Parallel Particle-Mesh Simulation using Nearest-Grid-Point (NGP)

Require: System *universe*, time step Δt, grid size N, spatial extent R, number of threads T

1: telemetry $\leftarrow \emptyset$
2: telemetry.append(initial positions)
3: boundaries $\leftarrow [-R, R]$
4: Compute cell size $h \leftarrow \frac{2R}{N}$
5: Initialize mass density grid $M[N][N] \leftarrow 0$
6: Initialize potential grid $\Phi[N][N] \leftarrow 0$
7: Initialize FFTW input/output arrays and plans
8: **for** each time step $s = 1$ to S **do**
9:     Clear mass density grid M
10:     **Parallel for** each thread $t = 1$ to T
11:         Assign a chunk of bodies to thread t
12:     **for** each body b assigned to thread t **do**
13:         $i \leftarrow \left\lfloor \frac{x_b - \min_x}{h} \right\rfloor$
14:         $j \leftarrow \left\lfloor \frac{y_b - \min_y}{h} \right\rfloor$
15:         **if** (i, j) in bounds **then**
16:             **Lock** $M[i][j]$
17:             $M[i][j] \leftarrow M[i][j] + \text{body.mass}$
18:             **Unlock** $M[i][j]$
19:         **end if**
20:     **end for**
21:     **End parallel for**
22:     Copy mass grid to FFTW input array
23:     Compute FFT of mass density using forward FFT
24:     **for** each (i, j) in frequency domain **do**
25:         Compute wave numbers (k_x, k_y)
26:         $k^2 \leftarrow k_x^2 + k_y^2$
27:         **if** $k^2 > 0$ **then**
28:             Multiply by $-\frac{G}{k^2}$ in frequency domain
29:         **else**
30:             Set value to zero
31:         **end if**
32:     **end for**
33:     Compute inverse FFT to obtain gravitational potential
34:     Normalize result and store in potential grid Φ
35:     **Parallel for** each thread $t = 1$ to T
36:         Assign a chunk of bodies to thread t
37:     **for** each body b assigned to thread t **do**
38:         $i \leftarrow \left\lfloor \frac{x_b - \min_x}{h} \right\rfloor$
39:         $j \leftarrow \left\lfloor \frac{y_b - \min_y}{h} \right\rfloor$
40:         **if** (i, j) valid and not at boundary **then**
41:             Compute force via central difference:
42:             $\vec{a}_b \leftarrow - \left(\frac{\Phi[i+1][j] - \Phi[i-1][j]}{2h}, \frac{\Phi[i][j+1] - \Phi[i][j-1]}{2h} \right)$
43:         **else**
44:             $\vec{a}_b \leftarrow (0, 0)$
45:         **end if**
46:     **end for**
47:     **End parallel for**
48:     **Parallel for** each thread $t = 1$ to T
49:         Assign a chunk of bodies $t = 1$ to T
50:     **for** each body b assigned to thread t **do**
51:         $\vec{v}_b \leftarrow \vec{v}_b + \vec{a}_b \Delta t$
52:         $\vec{r}_b \leftarrow \vec{r}_b + \vec{v}_b \Delta t$
53:     **end for**
54:     **End parallel for**
55:     telemetry.append(current positions)
56: **end for**
57: Cleanup FFTW plans and memory

4.1 Time performance of the particle-mesh simulation (sequential and parallel

The simulation where done on MacBook air with M2 chip containing 8 cores. All simulations where done with grid size $= 10$, a spatial extent R of 10000 and time increment $dt = 0.2$. Additionally, each simulation includes asteroids with high radius for their orbit (between 5 and 3000 as radius) and the rest being small asteroids with smaller radius for their orbit (between 0.5 and 5 as radius). The reason why I included big differences in radius is to allow between uniformity in the space domain.

Parallel computation was done through parallelizing via the bodies. Each body's mass was assigned in the grid via multiple threads and the computation of their acceleration was also done through parallelization.

The Fast Fourier transforms as well as the computation of the potential was not done in parallel since, assigning a thread would have been very inefficient if multiple bodies where assigned on the same grid. Indeed, locking the cells of each grid to avoid race conditions would have given a performance similar to a sequential one. In order to parallelize the computation of the gravitational potential, one needs smaller grids, and therefore divide the space into a higher value of grids. This would however mean higher number of "low-density" grids (grids with no or few bodies inside them) and therefore additional inefficient computation.

However, one can still implement Fast-Fourier transform in parallel within the library fftw3. Unfortunately, the performance observed was suboptimal.

Time performance of the sequential particle-mesh simulation

Number of Bodies	Execution Time (ms)
4	52
5	54
50	109
100	157
1,000	941
2,000	3,013
5,000	7,331
10,000	14,053

Table 1: Execution time of the particle-mesh simulation for varying numbers of bodies (grid size $= 10$)

Time performance of the parallel particle-mesh simulation

Number of Bodies	Threads $= 5$	Threads $= 7$	Threads $= 10$
4	1,661	–	–
5	2,209	2,848	4,058
50	2,295	2,906	3,263
100	2,298	3,034	3,165
1,000	3,436	3,800	5,006
2,000	5,075	5,865	6,595
5,000	7,576	6,244	6,757
10,000	18,230	12,306	13,881

Table 2: Execution time (in milliseconds) for varying numbers of bodies and thread counts (grid size $= 10$).

Looking at performance we see that until 2000 bodies, sequential time-performance is better than for the parallel implementation. However, for higher number of bodies, the parallel implementation with 7 threads displays better performance. 7 threads is particularly optimal as it is very close to the number of cores in the machine on which we simulate (8 cores), allowing optimal use of the CPU.

5 Barnes–Hut Algorithm

The Barnes–Hut algorithm reduces the complexity of the classical $O(N^2)$ N-body force computation to approximately $O(N \log N)$ by hierarchically clustering distant bodies. Our implementation follows these main stages:

Algorithm 4 Barnes–Hut Simulation Outline

Require: System of bodies with masses m_i, positions $\vec{r}_i$, velocities $\vec{v}_i$
Require: Time step Δt, number of steps N, opening angle θ
1: telemetry $\leftarrow \emptyset$
2: telemetry.append(initial positions)
3: **for** step $\leftarrow 0$ **to** $N - 1$ **do**
4:      $[\text{minB}, \text{maxB}] \leftarrow$ computeBounds(universe)
5:      root $\leftarrow$ createRootNode($[\text{minB}, \text{maxB}]$)
6:      **for** each body b_i **do**
7:          insertBody(root, b_i)
8:      **end for**
9:      computeMassDistribution(root)
10:      **for** each body b_i **do**
11:          $\vec{a}_i \leftarrow$ forceOnBody(b_i, root, θ)$/m_i$
12:      **end for**
13:      updateBodies(universe, Δt)
14:      freeQuadTree(root)
15:      telemetry.append(current positions)
16: **end for**
Ensure: Position history for all bodies stored in telemetry

---

1. **Compute Simulation Bounds.** We first call
   - `computeBounds(const System&)` to find an axis-aligned square enclosing all bodies (with a small padding).

2. **Quadtree Construction.** We represent space by a pointer-based quadtree of `QuadNode` objects:
   - `createRootNode(const Bounds&)` allocates the root covering the full region.
   - `insertBody(QuadNode *, Body&)` recursively subdivides nodes so that each leaf contains at most one body.

3. **Mass–Center Distribution.** A post-order traversal aggregates mass and center-of-mass at every internal node:
   - `computeMassDistribution(QuadNode *)` computes `totalMass` and `centerOfMass`.

4. **Force Computation.** For each body b_i, we traverse the tree and apply the opening-angle criterion θ:
   - `forceOnBody(const Body&, QuadNode, double theta)` approximates distant clusters as single masses, or recurses into children when closer.

5. **Parallelization (Shared-Memory).** To exploit multicore CPUs, we compute forces in parallel:
   - `computeForcesParallel(System&, QuadNode, double theta)` spawns a pool of `std::thread`s, partitions the bodies into chunks, and each thread calls `forceOnBody` on its subset.
   - Accelerations are stored per-thread and then written back, avoiding fine-grained locks.

6. **Time Integration.** Once all accelerations are known, we update velocities and positions in one pass:
   - `updateBodies(System&, double dt)` applies

$$\mathbf{v} \leftarrow \mathbf{v} + \mathbf{a}\,\Delta t, \quad \mathbf{x} \leftarrow \mathbf{x} + \mathbf{v}\,\Delta t.$$

7. **Cleanup.** The quadtree is deallocated to avoid memory leaks:

- `freeQuadTree(QuadNode *)` recursively `delete`s all nodes.

Concurrency Aspects

- **Multithreading in C++:** we illustrate basic shared-memory concurrency by partitioning the force computation across threads. This uses standard `std::thread` and per-thread buffers, avoiding complex locking.

- **Concurrent Data Structures:** insertion into the quadtree could be parallelized by feeding bodies into a thread-safe queue; our current version remains serial for clarity but is structured to allow a concurrent `insertBody` using a mutex per node or a lock-free pointer array.

- **GPU / PRAM Illustration (Optional):** under `USE_CUDA`, we provide `simulateBruteForceGPU(System&,double)` which launches an $O(N^2)$ CUDA kernel. Each GPU thread computes one body's net force, demonstrating the PRAM model in practice.

- **Correctness & Testing:** we compare parallel results to a single-threaded reference on small N, and use tools like ThreadSanitizer to detect data races. Telemetry outputs (positions over time) are also verified for physical invariants (e.g. center-of-mass motion).